


# Deploying Applications on Kubernetes — A Step-by-Step Practical Guide

**I**ntroduction: Kubernetes has become the de facto standard for container orchestration, offering scalability, fault tolerance, and ease of management for applications. In this step-by-step guide, we will walk you through the process of deploying applications on Kubernetes. Whether you're a beginner or an experienced developer, this guide will provide you with the necessary instructions to get your applications up and running on Kubernetes.



## Prerequisites

To begin, make sure you have the following prerequisites in place:

1. **A working Kubernetes cluster:** You can set up a cluster using tools like Minikube, Docker Desktop, or a cloud provider such as AWS, Azure, or Google Cloud.
2. **Kubectl:** The Kubernetes command-line tool (kubectl) should be installed and configured to communicate with your cluster.

3. **Docker:** Applications to be deployed on Kubernetes should be containerized using Docker. Ensure that Docker is installed and running on your local machine.

## Step 1: Create a Kubernetes Deployment

1. **Write a Dockerfile:** Create a Dockerfile for your application, specifying the necessary dependencies and instructions for building the container image.
2. **Build the Docker image:** Use the Dockerfile to build the container image for your application. Run the following command in the terminal:

```
docker build -t your-image-name:tag .
```

3. **Push the image to a registry:** If you're using a private registry, push the image to the registry using the following command:

```
docker push your-registry/your-image-name:tag
```

4. **Define a Deployment manifest:** Create a deployment YAML file (e.g., `deployment.yaml`) to describe the desired state of your application. Specify the image name, replica count, ports, and any other required configuration.

5. **Apply the Deployment:** Apply the deployment manifest using kubectl:

```
kubectl apply -f deployment.yaml
```

6. **Verify the Deployment:** Check the status of the deployment using the following command:

```
kubectl get deployments
```

## Step 2: Expose the Application

1. **Create a Service manifest:** To expose your application, create a service YAML file (e.g., `service.yaml`) that defines the networking rules for accessing the application.
2. **Specify the service type:** Choose the appropriate service type based on your requirements. This could be `NodePort`, `LoadBalancer`, or `ClusterIP`.
3. **Apply the Service:** Apply the service manifest using kubectl:

```
kubectl apply -f service.yaml
```

4. **Verify the Service:** Check the status of the service using the following command:

```
kubectl get services
```

### Step 3: Access the Application

1. Access the application locally: If you're using a **NodePort** service, find the port assigned to the service using **kubectl get services**. Access the application in your browser using the cluster IP and the assigned port.
2. Access the application externally: If you're using a **LoadBalancer** service, wait for the external IP to be assigned. Retrieve the external IP using **kubectl get services** and access the application using the IP.

**Conclusion:** Congratulations! You have successfully deployed your application on Kubernetes. By following this step-by-step guide, you learned how to create a Kubernetes deployment, expose your application through a service, and access it both locally and externally. Kubernetes offers a powerful platform for deploying and managing containerized applications at scale, and mastering its deployment process is essential for modern application development.

**Note:** Remember to clean up any resources created during this guide to avoid unnecessary costs or resource consumption. You can delete the deployment and service using **kubectl delete** with the appropriate YAML files.